

# TDT4165 Scala 101

Waclaw Kusnierczyk | [waclaw.kusnierczyk@gmail.com](mailto:waclaw.kusnierczyk@gmail.com)

# Agenda

1. What is Scala?
2. Hello :)
3. Execute!
4. OOPs
5. Pattern matching (or not)
6. Generics
7. Traits

# What is Scala?

Scala is a **scalable**, general-purpose programming language that

- is **compiled** (to Java bytecode), but has a strong **scripting** feel;
- is **statically typed**, but types can often be **inferred** by the compiler;
- has **object-oriented** programming features;
- has **functional** programming features;
- supports **operator overloading** and many other features absent from Java;
- and **much more**.

# Agenda

1. What is Scala?
2. Hello :)
3. Execute!
4. OOPs
5. Pattern matching (or not)
6. Generics
7. Traits

# Hello :)

Consider a simple `hello world` program in Java:

```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello, World");  
    }  
}
```

To make code executable, we need

- a class with a `public static void (String[])` method.

To run the program, we need to

- **compile** the source: `javac <source file name>;`
- **execute** the binary: `java <class file name>.`

# Hello :)

We can almost figuratively translate this to Scala:

```
object HelloWorld {  
  def main(args: Array[String]): Unit = {  
    println("Hello, World")  
  }  
}
```

# Hello :)

We can almost figuratively translate this to Scala:

```
object HelloWorld {  
  def main(args: Array[String]): Unit = {  
    println("Hello, World")  
  }  
}
```

## Note:

- `object` is a singleton instance **without** a class having ever been defined;
- `Array[String]` and `Unit` declare types **after** the declared objects;
- stuff is **public** unless explicitly **private** or **protected**.

# Hello :)

We can almost figuratively translate this to Scala:

```
object HelloWorld {  
  def main(args: Array[String]): Unit = {  
    println("Hello, World")  
  }  
}
```

## Note:

- `object` is a singleton instance **without** a class having ever been defined;
- `Array[String]` and `Unit` declare types **after** the declared objects;
- stuff is **public** unless explicitly **private** or **protected**.

Question: Where does `println` come from?



# Hello :)

Why be verbose? In Scala, we can avoid some of the syntactic noise:

```
object HelloWorld {  
  def main(args: Array[String]) =  
    println("Hello, World")  
}
```

# Hello :)

Why be verbose? In Scala, we can avoid some of the syntactic noise:

```
object HelloWorld {  
  def main(args: Array[String]) =  
    println("Hello, World")  
}
```

## Note:

- `main` calls `println` that does not return any value;
- hence, `main` does not return any value;
- the return type of `main` is thus inferred to be `Unit` (void in Java speak), so no need to be explicit;
- calling `main` on any `Array[String]` argument is equivalent to calling `println` on `"Hello, World!"`.

# Hello :)

We don't even need to explicitly implement `main`:

```
object HelloWorld extends App {  
  println("Hello, World!")  
}
```

# Hello :)

We don't even need to explicitly implement `main`:

```
object HelloWorld extends App {  
  println("Hello, World!")  
}
```

## Note:

- we `extend` the *trait* `App`: that's **inheritance** in action;
- `App` wraps our code in a ready-to-go `main`.

More on traits later.

Question: How do we run Scala code?

# Agenda

1. What is Scala?
2. Hello :)
3. Execute!
4. OOPs
5. Pattern matching (or not)
6. Generics
7. Traits

# Execute!

There are several ways to run Scala code:

- **compile** the source, **execute** the binary (as in Java):

```
$ scalac hello.scala # produces a binary
$ scala HelloWorld   # executes the binary
```

- **run** code directly from the source:

```
$ scala hello.scala # executes the source
```

## Note:

- the object `HelloWorld` does **not** have to be defined in a file named `HelloWorld.scala`.

# Execute!

There are several ways to run Scala code:

- run the code as a **script**:

```
$ ./hello.sh          # if hello.sh is executable
$ bash hello.sh       # otherwise
```

where `hello.sh` may be of two forms:

```
#!/bin/bash
exec scala "$0" "$@"
#!
object HelloWorld extends App {
  println("Hello, World!")
}
HelloWorld.main(null)
```

```
#!/bin/bash
exec scala "$0" "$@"
#!
println("Hello, World!")
```

# Execute!

There are several ways to run Scala code:

- start **interactive** Scala shell (just type `scala`), then type in the code:

```
scala> object HelloWorld {  
  | def main(args: Array[String]) = println("Hello, World!") }  
defined module HelloWorld  
scala> HelloWorld.main(null)  
Hello, World!  
scala> println("Hello, World!")  
Hello, World!
```



# Execute!

There are several ways to run Scala code:

- use **SBT**, the Scala build tool:

```
$ sbt          # choose action interactively in sbt shell
$ sbt run      # have sbt run code and exit
```

## Note:

- if multiple executable units are found (objects extending `App` or implementing `main`), you will be given a list to choose from:

```
Multiple main classes detected, select one to run:
[1] HelloWorld
[2] tdt4165.HelloWorld
Enter number:
```

# Agenda

1. What is Scala?
2. Hello :)
3. Execute!
4. OOPs
5. Pattern matching (or not)
6. Generics
7. Traits

# Object-oriented Scala

Use classes to specify the behaviour of a group of similar objects:

```
class Point(x: Int, y: Int) {  
  def show() = println(s"Point($x, $y)") }  
}
```

## Note:

- there is **no explicit constructor** method;
- parameters of the class (x, y) are accessible from within the class.

# Object-oriented Scala

Use classes to specify the behaviour of a group of similar objects:

```
class Point(x: Int, y: Int) {  
  def show() = println(s"Point($x, $y)") }  
}
```

Note:

- there is **no explicit constructor** method;
- parameters of the class (x, y) are accessible from within the class.

Question: What does `s"Point($x, $y)"` stand for?

# Object-oriented Scala

Instantiate classes with `new`:

```
val point = new Point(0, 0)
point.show()
```

## Note:

- `val` declares a new *value* (an **immutable variable**, so to speak);
- `new Point(0, 0)` returns an object of type `Point`;
- hence, the compiler can **infer** the type of `point`;
- hence, we do not need to; but we might:

```
val point: Point = new Point(0, 0)
point.show()
```

# Object-oriented Scala

Instantiate classes with `new`:

```
val point = new Point(0, 0)
point.show()
```

## Note:

- `val` declares a new *value* (an **immutable variable**, so to speak);
- `new Point(0, 0)` returns an object of type `Point`;
- hence, the compiler can **infer** the type of `point`;
- hence, we do not need to; but we might:

```
val point: Point = new Point(0, 0)
point.show()
```

**Question:** Why would we?

# Object-oriented Scala

The values `x` and `y` are visible exclusively inside the class `Point`:

```
val point = new Point(0, 0)
println(s"point = Point(${point.x}, ${point.y})") // x and y not found
```

# Object-oriented Scala

The values `x` and `y` are visible exclusively inside the class `Point`:

```
val point = new Point(0, 0)
println(s"point = Point(${point.x}, ${point.y})") // x and y not found
```

## Questions:

- Why not `s"point = Point(${point.x}, ${point.y})"`?
- Can we modify `x` and `y` (inside `Point`)?



# Object-oriented Scala

We can 'promote' class parameters to become true instance fields:

```
class Point(val x: Int, val y: Int) {  
  def show() = println(s"Point($x, $y)") }
```

Fields are (by default) public:

```
val point = new Point(0, 0)  
println(point.x)           // ok
```

# Object-oriented Scala

We can 'promote' class parameters to become true instance fields:

```
class Point(val x: Int, val y: Int) {  
  def show() = println(s"Point($x, $y)") }  
}
```

Fields are (by default) public:

```
val point = new Point(0, 0)  
println(point.x)           // ok
```

Question: Are the fields modifiable or mutable?

# Object-oriented Scala

We can 'promote' class parameters to become true instance fields:

```
class Point(val x: Int, val y: Int) {  
  def show() = println(s"Point($x, $y)") }  
}
```

Fields are (by default) public:

```
val point = new Point(0, 0)  
println(point.x)           // ok
```

Question: Are the fields modifiable or mutable? No, they aren't:

```
point.x = 1                 // error
```

# Object-oriented Scala

Let's augment the class to enable updating the coordinates:

```
class Point(var x: Int, var y: Int) {  
  def move(dx: Int, dy: Int) = {  
    x = x + dx  
    y = y + dy }  
  ... }  
val point = new Point(1, 1)
```

# Object-oriented Scala

Let's augment the class to enable updating the coordinates:

```
class Point(var x: Int, var y: Int) {  
  def move(dx: Int, dy: Int) = {  
    x = x + dx  
    y = y + dy }  
  ... }  
val point = new Point(1, 1)
```

## Questions:

- If `point` is a `val` (immutable variable), can `move` change its coordinates?
- Can the fields `x` and `y` be modified from outside of the class?

# Object-oriented Scala

Let's augment the class to enable updating the coordinates:

```
class Point(var x: Int, var y: Int) {  
  def move(dx: Int, dy: Int) = {  
    x = x + dx  
    y = y + dy }  
  ... }  
val point = new Point(1, 1)
```

## Questions:

- If `point` is a `val` (immutable variable), can `move` change its coordinates?
- Can the fields `x` and `y` be modified from outside of the class? Yes!

```
point.move(1, -1)    // Point(2, 0)  
point.y = 2          // Point(2, 2)
```

# Agenda

1. What is Scala?
2. Hello :)
3. Execute!
4. OOPs
5. Pattern matching (or not)
6. Generics
7. Traits

# Pattern-matching

We'd like to be able to compare a points for equality with other objects. Easy:

```
class Point(val x: Int, val y: Int) {  
  override def equals(any: Any): Boolean = {  
    if (!any.isInstanceOf[Point])  
      return false  
    val that = any.asInstanceOf[Point]  
    return that.x == this.x && that.y == this.y }  
  ... }
```

## Note:

- if you don't intend to modify fields, make them immutable (`vals`);
- the method `equals` is inherited from `Any` (Scala's `Object`);
- any **method overrides** must be explicitly acknowledged.



# Pattern-matching

We'd like to be able to compare a points for equality with other objects. Easy:

```
class Point(val x: Int, val y: Int) {  
  override def equals(any: Any): Boolean = {  
    if (!any.isInstanceOf[Point])  
      return false  
    val that = any.asInstanceOf[Point]  
    return that.x == this.x && that.y == this.y }  
  ... }
```

## Note:

- if you don't intend to modify fields, make them immutable (`vals`);
- the method `equals` is inherited from `Any` (Scala's `Object`);
- any **method overrides** must be explicitly acknowledged.

Question: This is kinda ugly. Can we do better?

# Pattern-matching

If we turn `Point` into a **case class**, we can use **pattern-matching**:

```
case class Point(x: Int, y: Int) {  
  override def equals(any: Any) = any match {  
    case Point(xx, yy) => xx == x && yy == y  
    case _ => false }  
  ... }  
}
```

## Note:

- in a `case` class, the parameters are `val` by default (accessible, immutable);
- `match` performs **pattern matching** between an object and one or more `case` patterns;
- a `case` pattern compares objects **structurally**, **capturing** variables (fields, etc.) as appropriate.
- the pattern `_` is a **catch-all** pattern that provides a default result.

# Pattern-matching

We can now easily compare a point to other objects:

```
val point = Point(0, 0)
point.equals(Point(1, 1)) // nope
point.equals(Point(0, 0)) // yepp
point.equals("Point(0, 0)") // nope!
```

## Note:

- a case class does not need `new` to construct instances.

# Pattern-matching

We can now easily compare a point to other objects:

```
val point = Point(0, 0)
point.equals(Point(1, 1)) // nope
point.equals(Point(0, 0)) // yepp
point.equals("Point(0, 0)") // nope!
```

## Note:

- a case class does not need `new` to construct instances.

# Agenda

1. What is Scala?
2. Hello :)
3. Execute!
4. OOPs
5. Pattern matching (or not)
6. Generics
7. Traits

# Generics

We could use different types of values for point coordinates:

```
case class Point[T](x: T, y: T) {  
  override def equals(any: Any) = any match {  
    case Point(xx: T, yy: T) => xx == x && yy == y  
    case Point(xx, yy) => // do some magic  
    case _ => false }  
  ... }
```

## Note:

- `Point[T]` is a **generic class**, where `T` is a type variable;
- both arguments are of type `T`;
- the pattern `case Point(xx: T, yy: T)` matches all and only `T`-points.

# Generics

We could use different types of values for point coordinates:

```
case class Point[T](x: T, y: T) {  
  override def equals(any: Any) = any match {  
    case Point(xx: T, yy: T) => xx == x && yy == y  
    case Point(xx, yy) => // do some magic  
    case _ => false }  
  ... }
```

Note:

- `Point[T]` is a **generic class**, where `T` is a type variable;
- both arguments are of type `T`;
- the pattern `case Point(xx: T, yy: T)` matches all and only `T`-points.

Question: What is the outcome of:

```
Point(0, 1/2).equals(Point(0, 0.5))
```

# Agenda

1. What is Scala?
2. Hello :)
3. Execute!
4. OOPs
5. Pattern matching (or not)
6. Generics
7. Traits



# Traits

Suppose we are allergic to `!` and prefer `a unequals b` to `!a.equals(b)`.

```
case class Point[T](x: T, y: T) {  
  override def equals(any: Any) = any match {  
    case Point(xx: T, yy: T) => xx == x && yy == y  
    case _ => false }  
  def unequals(any: Any) = !equal(any)  
  ... }  
}
```

We can now write:

```
Point(0, 0) unequals Point(1, 1) // true
```

instead of:

```
!Point(0, 0).equals(Point(1, 1)) // true
```

(Yeah, **silly**. But note the elegant **infix notation**!)

# Traits

We don't want to implement both `equals` and `unequals` for each new class.

Let's implement a reusable **trait**:

```
trait Unequal {  
  def unequals(any: Any) = !equals(any) }  
case class Point[T](x: T, y: T) extends Unequal {  
  ... }
```

## Note:

- `Unequals` is not a class, and cannot be instantiated;
- the **trait provides** a method that can be immediately used by any class that extends it;
- it refers to `equals`, a method from `Any` that every class either inherits or overrides.

# Traits

We don't want to implement both `equals` and `unequals` for each new class.

Let's implement a reusable **trait**:

```
trait Unequal {  
  def unequals(any: Any) = !equals(any) }  
case class Point[T](x: T, y: T) extends Unequal {  
  ... }
```

## Note:

- `Unequals` is not a class, and cannot be instantiated;
- the **trait provides** a method that can be immediately used by any class that extends it;
- it refers to `equals`, a method from `Any` that every class either inherits or overrides.

**Question:** How novel is this?

# Traits

We don't want to implement both `equals` and `unequals` for each new class.

Let's implement a reusable **trait**:

```
trait Unequal {  
  def unequals(any: Any) = !equals(any) }  
case class Point[T](x: T, y: T) extends Unequal {  
  ... }
```

## Note:

- `Unequals` is not a class, and cannot be instantiated;
- the **trait provides** a method that can be immediately used by any class that extends it;
- it refers to `equals`, a method from `Any` that every class either inherits or overrides.

**Question:** How novel is this? (Think interfaces with default method implementations in Java.)

# TDT4165: Concurrency in Scala, Part I

October 28, 2015

Slide 1 of 32

# Processes and Threads

## A **process**

- is an instance of a computer program that is being executed;
- is largely **self-contained**: has its own, **isolated** share of memory and other computational resources;
- cannot directly access the memory of other processes;
- cannot simultaneously use the same resources as other processes.

# Processes and Threads

## A **process**

- is an instance of a computer program that is being executed;
- is largely **self-contained**: has its own, **isolated** share of memory and other computational resources;
- cannot directly access the memory of other processes;
- cannot simultaneously use the same resources as other processes.

## A **thread**

- is a **unit of execution** within a process;
- has its own program stack and program counter;
- **shares memory and resources** with other threads within the process.

# Processes and Threads

## A **process**

- is an instance of a computer program that is being executed;
- is largely **self-contained**: has its own, **isolated** share of memory and other computational resources;
- cannot directly access the memory of other processes;
- cannot simultaneously use the same resources as other processes.

## A **thread**

- is a **unit of execution** within a process;
- has its own program stack and program counter;
- **shares memory and resources** with other threads within the process.

**Note:** We will be concerned mostly with threads (i.e., **multithreaded** programs).



# Processes and Threads

For debugging, it will be useful to **identify** threads.

```
object ThreadName extends App {  
  val name = Thread.currentThread.getName  
  println(s"thread: $name") }  
}
```

## Note:

- for convenience, we **extend the trait** `App` to inherit `main` (see `scala.App`);
- `Thread` is a **Java** class (see `java.lang.Thread`);
- `currentThread` is a `static` method of `Thread` that *returns a reference to the currently executing thread object* (an instance of `Thread`);
- `getName` is an instance method of `Thread` that *returns this thread's name*;

# Processes and Threads

For debugging, it will be useful to **identify** threads.

```
object ThreadName extends App {  
  val name = Thread.currentThread.getName  
  println(s"thread: $name") }  
}
```

## Note:

- for convenience, we **extend the trait** `App` to inherit `main` (see `scala.App`);
- `Thread` is a **Java** class (see `java.lang.Thread`);
- `currentThread` is a static method of `Thread` that *returns a reference to the currently executing thread object* (an instance of `Thread`);
- `getName` is an instance method of `Thread` that *returns this thread's name*;

## Question:

- Why `currentThread` and `getName`, and not `currentThread()` and `getName()`?

# Processes... Methods

Method definition syntax:

- `currentThread` and `getName` are **nullary** methods: they take **no arguments** (and **return** a value);
- in Scala, nullary methods can be defined with or without (empty) parentheses;
- methods defined with empty parentheses can be called with or without them.

```
def f1 = "f1"    // def f1: String = ...  
def f2() = "f2"  // def f2(): String = ...
```

# Processes... Methods

Method definition syntax:

- `currentThread` and `getName` are **nullary** methods: they take **no arguments** (and **return** a value);
- in Scala, nullary methods can be defined with or without (empty) parentheses;
- methods defined with empty parentheses can be called with or without them.

```
def f1 = "f1"    // def f1: String = ...  
def f2() = "f2"  // def f2(): String = ...
```

Note:

- both `f1` and `f2` are nullary methods returning a value of type `String`;
- their definitions are **roughly equivalent**, but
- there are differences in how they can be **called**.

# Processes... Methods

How a nullary method should be called depends on its definition.

```
def f1 = "f1"    // def f1: String = ...  
def f2() = "f2"  // def f2(): String = ...
```

```
println(f1)      // (1)  
println(f1())    // (2)  
println(f1(0))   // (3)  
println(f2)      // (4)  
println(f2())    // (5)  
println(f2(0))   // (6)
```

## Questions:

- Are any/all of the calls legal? If legal, what is the output?

# Processes... Methods

How a nullary method should be called depends on its definition.

```
def f1 = "f1"    // def f1: String = ...  
def f2() = "f2"  // def f2(): String = ...
```

```
println(f1)      // (1)  
println(f1())    // (2)  
println(f1(0))   // (3)  
println(f2)      // (4)  
println(f2())    // (5)  
println(f2(0))   // (6)
```

## Questions:

- Are any/all of the calls legal? If legal, what is the output?

## Answer:

- 1: f1 | 2: **incidentally illegal** | 3: f, **incidentally legal**
- 4: f2 | 5: f2 | 6: illegal.

# Processes... Methods

It is **customary** to use `()` in both definitions and calls to methods that **do not return** values (i.e., have `Unit` return type).

```
def p1 = println("p1")    // def p1: Unit = ...  
def p2() = println("p2")  // def p2(): Unit = ..., preferred
```

```
p1      // (1)  
p1()    // (2)  
p2      // (3)  
p2()    // (4)
```

## Question:

- Are any/all of the calls legal?

# Processes... Methods

It is **customary** to use `()` in both definitions and calls to methods that **do not return** values (i.e., have `Unit` return type).

```
def p1 = println("p1")    // def p1: Unit = ...  
def p2() = println("p2")  // def p2(): Unit = ..., preferred
```

```
p1      // (1)  
p1()    // (2)  
p2      // (3)  
p2()    // (4)
```

## Question:

- Are any/all of the calls legal?

## Answer:

- 1: legal | 2: illegal
- 3: legal | 4: legal, **preferred**



# Processes and Threads

A convenience **library** method to print messages with thread id prefixed:

```
package cis.utils
object ThreadUtils {
  def log(message: String): Unit =
    println(s"[${Thread.currentThread.getName}] $message") }
```

**Note:** Unlike in Java, in Scala

- a class (or object) does **not** have to reside in a file named exactly after the class;
- the `package` declaration does **not** have to reflect the directory structure above the source file;
- a source file may include **more than one** class (or object).

# Processes... Project Structure

To ease building, organize code into a **standard structure**.

```
$ tree
|-- src
|   |-- main
|       |-- resources
|       |-- scala
|           |-- examples
|           |   |-- ThreadsDemo.scala # code with main
|           |-- utils
|           |   |-- Threads.scala     # library code
|-- build.sbt
```

Include a rudimentary **build file** `build.sbt` at the project root:

```
scalaVersion := "2.11.7"
```

# Processes and Threads

We can now conveniently print diagnostic messages **prefixed** with a thread id:

```
package cis.example
import cis.utils.ThreadUtils._
object MainThreadDemo extends App {
  log(s"reetings from $this!") }
```

# Processes and Threads

We can now conveniently print diagnostic messages **prefixed** with a thread id:

```
package cis.example
import cis.utils.ThreadUtils._
object MainThreadDemo extends App {
  log(s"reetings from $this!") }
```

## Note:

- the blank in `ThreadUtils._` matches anything, hence
- we are importing the **whole content** of the `ThreadUtils` object;
- equivalent to `import static cis.utils.ThreadUtils.*` in Java;
- we can use `log` as if it were a locally defined method.
- some prefer explicit imports, e.g., `import cis.utils.ThreadUtils.log`

# Processes and Threads

We can now conveniently print diagnostic messages **prefixed** with a thread id:

```
package cis.example
import cis.utils.ThreadUtils._
object MainThreadDemo extends App {
  log(s"greetings from $this!") }
```

## Note:

- the blank in `ThreadUtils._` matches anything, hence
- we are importing the **whole content** of the `ThreadUtils` object;
- equivalent to `import static cis.utils.ThreadUtils.*` in Java;
- we can use `log` as if it were a locally defined method.
- some prefer explicit imports, e.g., `import cis.utils.ThreadUtils.log`

## Question:

- What does `$this` evaluate to within the message?

# Processes and Threads

Let us start a new thread:

```
object NewThreadDemo extends App {  
  val thread = new Thread {  
    override def run() = log(s"greetings from $this!") } }  
}
```

## Note:

- `log` is imported from `ThreadUtils`, as before;
- from now on, **we will skip** package and `import` lines where they are obvious;
- as with nullary functions, the **constructor** `Thread()` is called **without parentheses**;
- we create an instance of an anonymous class that extends `Thread`, overriding its `run` method.

# Processes and Threads

Let us start a new thread:

```
object NewThreadDemo extends App {  
  val thread = new Thread {  
    override def run() = log(s"greetings from $this!") } }  
}
```

## Note:

- `log` is imported from `ThreadUtils`, as before;
- from now on, **we will skip** package and `import` lines where they are obvious;
- as with nullary functions, the **constructor** `Thread()` is called **without parentheses**;
- we create an instance of an anonymous class that extends `Thread`, overriding its `run` method.

## Question:

- What will be the result of running this program?

# Processes and Threads

To avoid repetitive typing, let's add the thread creating code to our library:

```
object ThreadUtils {  
  def thread(body: => Unit): Thread =  
    new Thread {  
      override def run() = body }  
  ... }  
}
```



# Processes and Threads

To avoid repetitive typing, let's add the thread creating code to our library:

```
object ThreadUtils {  
  def thread(body: => Unit): Thread =  
    new Thread {  
      override def run() = body }  
  ... }  
}
```

**Note:** We update `ThreadUtils` by adding

- the method `thread` that
- takes as an argument an **unevaluated statement** `body`, and
- returns a new instance (a thread) of an anonymous class that
- extends `Thread` implementing `run` that
- executes `body`.

# Processes and Threads

To avoid repetitive typing, let's add the thread creating code to our library:

```
object ThreadUtils {  
  def thread(body: => Unit): Thread =  
    new Thread {  
      override def run() = body  
    }  
  ...  
}
```

**Note:** We update ThreadUtils by adding

- the method thread that
- takes as an argument an **unevaluated statement** body, and
- returns a new instance (a thread) of an anonymous class that
- extends Thread implementing run that
- executes body.

## Question:

- What exactly is body?

# Processes... Function Objects

Scala allows us to create **anonymous function** (and **procedure**) objects.

```
val double = new Function1[Int, Int] {  
  def apply(int: Int): Int = 2 * int  
}
```

**Note:** `double`

- is a **unary** function from `Int` to `Int` that
- is an instance of an **anonymous class extending** the trait `Function1`, which
- **implements** `apply` declared in `Function1`.

# Processes... Function Objects

Scala allows us to create **anonymous function** (and **procedure**) objects.

```
val double = new Function1[Int, Int] {  
  def apply(int: Int): Int = 2 * int }
```

**Note:** double

- is a **unary** function from Int to Int that
- is an instance of an **anonymous class extending** the trait Function1, which
- **implements** apply declared in Function1.

## Questions:

- Don't we have to **explicitly override** parent methods?
- Why don't we need (though we can) override def apply here?

# Processes... Function Objects

Being lazy programmers, we appreciate more **compact syntax**.

```
val double = (int: Int) => 2 * int
```

**Note:** As previously,

- `double` is an anonymous unary function with
- **explicit** input type `Int`, and
- **implicit** output type `(Int)`.

# Processes... Function Objects

Being lazy programmers, we appreciate more **compact syntax**.

```
val double = (int: Int) => 2 * int
```

**Note:** As previously,

- `double` is an anonymous unary function with
- **explicit** input type `Int`, and
- **implicit** output type `(Int)`.

## Questions:

- Is the type of `double` itself explicit or implicit? What is `double`'s type?

# Processes... Function Objects

Being lazy programmers, we appreciate more **compact syntax**.

```
val double = (int: Int) => 2 * int
```

**Note:** As previously,

- `double` is an anonymous unary function with
- **explicit** input type `Int`, and
- **implicit** output type `(Int)`.

## Questions:

- Is the type of `double` itself explicit or implicit? What is `double`'s type?

Consider:

```
val double: Int => Int = (int: Int) => 2 * int
```

# Processes... Function Objects

Functions in Scala are **first-class objects**.

```
def twice[T] =  
  (func: (T => T)) =>  
    (input: T) => func(func(input))
```

## Note:

- `twice` is a T-generic method whose value is an anonymous function that
- **takes** as an argument a **function** `func` from T to T and
- **returns** a **function** that takes an `input` of type T and
- returns the result of `func` applied to the result of `func` applied to `input`.



# Processes... Function Objects

Functions in Scala are **first-class objects**.

```
def twice[T] =  
  (func: (T => T)) =>  
    (input: T) => func(func(input))
```

## Note:

- `twice` is a T-generic method whose value is an anonymous function that
- **takes** as an argument a **function** `func` from T to T and
- **returns** a **function** that takes an `input` of type T and
- returns the result of `func` applied to the result of `func` applied to `input`.

## Questions:

- What is the value of `twice(twice[Int])(double)(2)`?
- What is the type of `twice[T]`?

# Processes and Threads

**Wrapping** code to be evaluated by a thread into a nullary procedure or function (a **thunk**), we can:

- **delay** its evaluation until a call to the thunk is made;
- **pass** it in unevaluated form **as an argument** to a function.

```
object ThreadUtils {  
  def thread(body: () => Unit): Thread =  
    new Thread { override def run() = body() }  
  ... }
```

## Note:

- with `body: () => Unit ... run() = body()`, you **pass by value** a nullary procedure to be called with no arguments by the thread;
- with `body: => Unit ... run() = body`, you **pass by name** an unevaluated statement to be evaluated by the thread;
- `=> Unit` is a **syntactic extension** that is resolved to `() => Unit` at compile time. (Hint: try to overload `def thread` with both signatures.)

# Processes and Threads

Consider this example:

```
object NewThreadDemo extends App {  
  log("1")  
  val t = thread { log("2"); Thread.sleep(1000); log("3") }  
  log("4")  
  t.start()  
  log("5")  
  Thread.sleep(1000)  
  log("6")  
  t.join()  
  log("7") }
```

## Note:

- `Thread.sleep(1000)` makes the **calling** thread **pause for at least** 1000 ms;
- `t.join()` makes the **calling** thread (main) **wait until** thread t **completes**.

Questions: What can we say about the output?

# Processes and Threads

Consider this example:

```
object NewThreadDemo extends App {  
  log("1")  
  val t = thread { log("2"); Thread.sleep(1000); log("3") }  
  log("4")  
  t.start()  
  log("5")  
  Thread.sleep(1000)  
  log("6")  
  t.join()  
  log("7") }  
}
```

## Note:

- `Thread.sleep(1000)` makes the **calling** thread **pause for at least** 1000 ms;
- `t.join()` makes the **calling** thread (main) **wait until** thread t **completes**.

Questions: What can we say about the output? **Test it!**

# Processes and Threads

If a program is executed in **more than one thread**, then...

# Processes and Threads

If a program is executed in **more than one thread**, then...

- the overall behaviour will usually be **nondeterministic**;
- the execution may lead to **unexpected results**;
- the execution may lead to **incorrect results**;
- the execution may lead to **no results** at all;
- **testing** is difficult and tricky.

# Processes and Threads

If a program is executed in **more than one thread**, then...

- the overall behaviour will usually be **nondeterministic**;
- the execution may lead to **unexpected results**;
- the execution may lead to **incorrect results**;
- the execution may lead to **no results** at all;
- **testing** is difficult and tricky.

## How do you change a lightbulb in...

- Object-oriented programming? You don't. Instead, you tell the lamp to do it.
- Functional programming? You can't.
- Logical programming? You imply that it is changed.
- **Concurrent programming**? You take the lamp to a secure area where nobody else can try to change the lightbulb while you're changing it, and do an atomic swap of the old lightbulb with the new lightbulb.

# Processes and Threads

```
log("1")
val t = thread { log("2"); Thread.sleep(1000); log("3") }
log("4")
t.start()
log("5")
Thread.sleep(1000)
log("6")
t.join()
log("7")
```

In the `NewThreadDemo` example, we can say that the output

- **always** starts with 1 4;
- **always** ends with 7;
- **always** contains 2 before 3 and 5 before 6;
- **almost certainly** contains 2 before 6 and 5 before 3;
- **may** contain 2 before 5 or vice versa, and likewise with 3 and 6.



# Processes and Threads

Let us create another convenience method in `ThreadUtils`:

```
def start(body: => Unit, delay: Int = 0): Thread = {  
  val t = thread { Thread.sleep(delay); body }  
  t.start()  
  t }
```

## Note:

- `start` calls the previously defined `thread`;
- the body passed to `thread` contains a `Thread.sleep` call and the body passed to `start`;
- `delay` has a **default** value 0 and is **optional** in calls to `start`;
- a call to `start` will create a new thread that will execute `body` after a set `delay`, start the thread, and return it.

# Processes and Threads

Mutable state and concurrent execution **interact unpredictably**.

```
var state = 0
val threads = (1 to n).map(i => start({ state = (state + i) * i }, delay))
threads.foreach(_.join())
```

## Note:

- `1.to(n)` is a **sequence** of `Int`s from 1 to `n`, inclusive;
- `1 to n` is **infix** form of `1.to(n)`;
- `map` **applies** a unary function to each element of the sequence, and
- returns a **new sequence** of the results;
- `i => start(...)` takes an `Int` and returns a **started** `Thread` that
- sleeps for `delay` milliseconds, and then **destructively updates** `state` using `i`;
- `threads` is a **sequences of threads**;
- `foreach` **iterates** over `threads` to join each in turn;
- `_.join()` is **syntactic sugar** for `thread => thread.join()`.

# Processes and Threads

```
def test(n: Int, delay: Int = 0) = {  
  var state = 0  
  val threads = (1 to n).map(i => start({ state = (state + i) * i }, delay))  
  threads.foreach(_.join())  
  state  
}
```

**Question:** What is the value of `test(2)`? `test(3)`?

- The **order** of execution of multiple threads is **unpredictable**.
- For `test(2)`, we would expect

# Processes and Threads

```
def test(n: Int, delay: Int = 0) = {  
  var state = 0  
  val threads = (1 to n).map(i => start({ state = (state + i) * i }, delay))  
  threads.foreach(_.join())  
  state  
}
```

**Question:** What is the value of `test(2)`? `test(3)`?

- The **order** of execution of multiple threads is **unpredictable**.
- For `test(2)`, we would expect either  $((0+1)*1+2)*2 = 6$  or  $((0+2)*2+1)*1 = 5$ .
- For `test(3)`, we would expect 22, 23, 24, 27, or 28.

# Processes and Threads

```
def test(n: Int, delay: Int = 0) = {  
  var state = 0  
  val threads = (1 to n).map(i => start({ state = (state + i) * i }, delay))  
  threads.foreach(_.join())  
  state  
}
```

**Question:** What is the value of `test(2)`? `test(3)`?

- The **order** of execution of multiple threads is **unpredictable**.
- For `test(2)`, we would expect either  $((0+1)*1+2)*2 = 6$  or  $((0+2)*2+1)*1 = 5$ .
- For `test(3)`, we would expect 22, 23, 24, 27, or 28.

But (with sufficiently large `delay`) **we actually get:**

- 1, 4, 5, 6 (for `test(2)`);
- 1, 4, 5, 6, 9, 10, 12, 21, 22, 23, 24, 27, 28 (for `test(3)`), etc.

# Processes and Threads

```
def test(n: Int, delay: Int = 0) = {  
  var state = 0  
  val threads = (1 to n).map(i => start({ state = (state + i) * i }, delay))  
  threads.foreach(_.join())  
  state  
}
```

**Question:** What is the value of `test(2)`? `test(3)`?

- The **order** of execution of multiple threads is **unpredictable**.
- For `test(2)`, we would expect either  $((0+1)*1+2)*2 = 6$  or  $((0+2)*2+1)*1 = 5$ .
- For `test(3)`, we would expect 22, 23, 24, 27, or 28.

But (with sufficiently large `delay`) **we actually get**:

- 1, 4, 5, 6 (for `test(2)`);
- 1, 4, 5, 6, 9, 10, 12, 21, 22, 23, 24, 27, 28 (for `test(3)`), etc.

**Questions:** Why are all those **unexpected** results **smaller** than those expected?

# Processes and Threads

With concurrency and mutable state, predicting is not easy. Experiment!

```
def tests(ntests: Int = 100, nthreads: Int = 1, delay: Int = 0) =  
  (1 to ntests).par  
    .map(_ => test(nthreads, delay))  
    .groupBy(identity).mapValues(_.size)  
    .toSeq.seq.sortBy(_._1)
```

## Note:

- `par` **parallelizes** the sequence (further steps can be run concurrently);
- `map` **executes** test for each value in `(1 to ntests)`;
- `groupBy` **aggregates** the return values into a `Map` with unique results as keys and sequences of those as values;
- `mapValues` **calculates** the length of each sequence value in the map;
- `toSeq` **converts** the map into a sequence of (result, size) tuples;
- `seq` **de-parallelizes** the sequence;
- `sortBy` **reorders** the tuples by the first element.

# Processes and Threads

The update of `state` is not atomic:

```
state = (state + i) * i
```

is equivalent to

```
val temp1 = state
val temp2 = temp1 + i
val temp3 = temp2 * i
state = temp3
```

Note:

- `state` is the only **mutable, shared** state variable;
- `i`, `temp1`, `temp2`, and `temp3` are thread-local and immutable;
- threads may access `state` in **between read and write** operations.



# Processes and Threads

In `test(2)`, one possible sequence of events is

- thread 1 | `val temp1 = state` (`state = 0`, thus `temp1 = 0`)
- thread 2 | `val temp1 = state` (`state = 0`, thus `temp1 = 0`)
- thread 2 | `val temp2 = temp1 + i` (`i = 2`, `temp1 = 0`, thus `temp2 = 2`)
- thread 1 | `val temp2 = temp1 + i` (`i = 1`, `temp1 = 0`, thus `temp2 = 1`)
- thread 1 | `val temp3 = temp2 * i` (`i = 1`, `temp2 = 1`, thus `temp3 = 1`)
- thread 2 | `val temp3 = temp2 * i` (`i = 2`, `temp2 = 2`, thus `temp3 = 4`)
- thread 2 | `state = temp3` (`temp3 = 4`, thus `state = 4`)
- thread 1 | `state = temp3` (`temp3 = 1`, thus `state = 1`)

# Processes and Threads

In `test(2)`, one possible sequence of events is

- thread 1 | `val temp1 = state` (`state = 0`, thus `temp1 = 0`)
- thread 2 | `val temp1 = state` (`state = 0`, thus `temp1 = 0`)
- thread 2 | `val temp2 = temp1 + i` (`i = 2`, `temp1 = 0`, thus `temp2 = 2`)
- thread 1 | `val temp2 = temp1 + i` (`i = 1`, `temp1 = 0`, thus `temp2 = 1`)
- thread 1 | `val temp3 = temp2 * i` (`i = 1`, `temp2 = 1`, thus `temp3 = 1`)
- thread 2 | `val temp3 = temp2 * i` (`i = 2`, `temp2 = 2`, thus `temp3 = 4`)
- thread 2 | `state = temp3` (`temp3 = 4`, thus `state = 4`)
- thread 1 | `state = temp3` (`temp3 = 1`, thus `state = 1`)

OOPS!

# Synchronization

To prevent undesirable thread interactions, **protect shared, mutable** state.

```
def test(n: Int, delay: Int = 0) = {  
  var state = 0  
  val threads = (1 to n)  
    .map(i => start(this.synchronized { state = (state + i) * i }, delay)  
    ... }
```

## Note:

- `this` refers to the **enclosing object** (one that `test` is a method of);
- the `synchronized` statement wraps a block to be executed **atomically**;
- `synchronized` causes a thread to **obtain a lock** on `this`;
- the thread **keeps the lock** until ready with executing the block;
- when `this` is locked, **no other thread** can lock it; hence
- **only one thread** at a time can be executing the block;
- the block can be executed only as a **whole** (atomically).

# Synchronization

In `test(2)`, one possible sequence of events is

- thread 1 | **lock**
- thread 1 | `val temp1 = state` (`state = 0`, thus `temp1 = 0`)
- thread 2 | **cannot lock**
- thread 1 | `val temp2 = temp1 + i` (`i = 1`, `temp1 = 0`, thus `temp2 = 1`)
- thread 1 | `val temp3 = temp2 * i` (`i = 1`, `temp2 = 1`, thus `temp3 = 1`)
- thread 2 | **cannot lock**
- thread 1 | `state = temp3` (`temp3 = 1`, thus `state = 1`)
- thread 1 | **unlock**
- thread 2 | **lock**
- thread 2 | `val temp1 = state` (`state = 1`, thus `temp1 = 1`)
- thread 2 | `val temp2 = temp1 + i` (`i = 2`, `temp1 = 1`, thus `temp2 = 3`)
- thread 2 | `val temp3 = temp2 * i` (`i = 2`, `temp2 = 3`, thus `temp3 = 6`)
- thread 2 | `state = temp3` (`temp3 = 6`, thus `state = 6`)
- thread 2 | **unlock**

# Quiz!

Threads synchronizing on **the same object** may **execute different code**.

Question: What is the most likely value to be printed by this program?

# Quiz!

Threads synchronizing on **the same object** may **execute different code**.

**Question:** What is the most likely value to be printed by this program?

```
object Quiz extends App {  
  var state = 10  
  start(this.synchronized { state *= 10 }, 1000)  
  start(this.synchronized { state += 10 }, 1000)  
  println(s"$state") }
```

# Synchronization

Synchronization solves one problem, but introduces others.

Consider the classic problem of **dining philosophers**.

```
case class Fork()
case class Philo(left: Fork, right: Fork) {
  def eat() = left.synchronized { right.synchronized { log(s"$this") } } }
class Dinner(val n: Int) {
  val forks = (0 until n).map(_ => Fork())
  val philos = (0 until n).map(i => Philo(forks(i), forks((i+1)%n))) }
```

# Synchronization

Synchronization solves one problem, but introduces others.

Consider the classic problem of **dining philosophers**.

```
case class Fork()
case class Philo(left: Fork, right: Fork) {
  def eat() = left.synchronized { right.synchronized { log(s"$this") } } }
class Dinner(val n: Int) {
  val forks = (0 until n).map(_ => Fork())
  val philos = (0 until n).map(i => Philo(forks(i), forks((i+1)%n))) }
```

Question: What will happen if we throw a party?

```
new Dinner(10).philos.foreach(_.eat())
```



# Synchronization

Synchronization solves one problem, but introduces others.

Consider the classic problem of **dining philosophers**.

```
case class Fork()
case class Philo(left: Fork, right: Fork) {
  def eat() = left.synchronized { right.synchronized { log(s"$this") } } }
class Dinner(val n: Int) {
  val forks = (0 until n).map(_ => Fork())
  val philos = (0 until n).map(i => Philo(forks(i), forks((i+1)%n))) }
```

Question: What will happen if we throw a party?

```
new Dinner(10).philos.foreach(_.eat())
```

It's a **kølapp party**, no deadlock.

# Synchronization

How about a real party?

```
new Dinner(10).philos.par.foreach(_.eat())
```


## Note:

- the philosophers start (or attempt to start) eating simultaneously;
- each philosopher grabs the left fork;
- no philosopher can grab the right fork;
- the **execution halts** indefinitely.

# Synchronization

**Deadlock** is a dangerous consequence of too much synchronization.

- The dining philosophers has some trivial solutions.
- In general, chances for deadlock can be reduced by synchronizing over the **smallest possible blocks** of code and on the **most local objects** possible.
- Other ways to assure atomicity may help avoid deadlocks entirely.



# Compare and Swap

Consider the task of generating successive unique ids.

```
object UIDGenerator {  
  var uid = 0;  
  def getNext = {  
    uid += 1  
    uid } }  
}
```

## Note:

- This implementation is **flawed**.
- The bug is solved by wrapping the update `uid += 1` in a `synchronized` block.
- It may be more convenient to use a **compare-and-swap** (CAS) operation.

# Compare and Swap

```
import java.util.concurrent.atomic.AtomicInteger
import scala.annotation.tailrec
object UIDGenerator {
  val uid = new AtomicInteger(0)
  @tailrec def getNext: Int = {
    val current = uid.get
    val updated = current + 1
    if uid.compareAndSet(current, updated) return updated
    getNext }
}
```

## Note:

- an `AtomicInteger` is a **wrapper** around an `Int` that
- allows one to **replace** the content if one knows the current value
- `getNext` is a **recursive** (self-calling) method that gets the current value, calculates an update, and tries to swap them;
- `getNext` **fails** if another thread has modified `uid` since `current` was obtained;
- `@tailrec` states that `getNext` is **tail recursive**.

# TDT4165: Concurrency in Scala, Part II

November 4, 2015

Slide 1 of 21

# Lazy Values

A **lazy value** is a value that

- remains **uninitialized** until it is needed;
- is initialized the **first time** it is used,
- is initialized only **once**.

```
lazy val value = { /* some horribly involved computation */ }
```

## Note:

- lazy values are created with the keyword `lazy`;
- there are **no lazy variables**.

# Lazy Values

A **lazy value** is a value that

- remains **uninitialized** until it is needed;
- is initialized the **first time** it is used,
- is initialized only **once**.

```
lazy val value = { /* some horribly involved computation */ }
```

## Note:

- lazy values are created with the keyword `lazy`;
- there are **no lazy variables**.

With lazy values, potentially complex computations

- are **avoided** if their result is never used;
- may have **surprising results** in sequential execution context;
- may have **unpredictable results** in concurrent execution context.



# Lazy Values

In a sequential program, lazy values may lead to surprising, but predictable behaviour.

```
var square = 121
lazy val root = math.sqrt(square).toInt
var square = 0
log(s"square: $square, root: $root")
```

**Question:** What can we say about the output?

# Lazy Values

In a sequential program, lazy values may lead to surprising, but predictable behaviour.

```
var square = 121
lazy val root = math.sqrt(square).toInt
var square = 0
log(s"square: $square, root: $root")
```

**Question:** What can we say about the output?

**Note:**

- `root` is first used in the `println` statement;
- `root` is uninitialized until then;
- when `root` is initialized, `square` is already reassigned the value 0;
- hence, `root = 0`.

# Lazy Values

In a concurrent program, lazy values may lead to unpredictable behaviour.

```
var root = 0
lazy val square = root * root
val thread = (1 to 3)
  .map(i => start({ root = i; log(s"square: $square") }, 100))
```

Question: What can we say about the output?

# Lazy Values

In a concurrent program, lazy values may lead to unpredictable behaviour.

```
var root = 0
lazy val square = root * root
val thread = (1 to 3)
    .map(i => start({ root = i; log(s"square: $square") }, 100))
```

Question: What can we say about the output?

**Note**:

- three different threads are started (in addition to main);
- the **first thread** that reaches its `println` statement **forces** `square` to be initialized;
- the remaining threads see `square` **already initialized** and print the same value as the first thread.

# Lazy Values

**Streams** (aka *lazy lists*) are a prime use example for lazy values.

```
def ints(first: Int): Stream[Int] =  
  Stream.cons(first, ints(first + 1))  
def primes(ints: Stream[Int]): Stream[Int] =  
  Stream.cons(ints.head, primes(ints.tail filter (_ % ints.head != 0)))  
def primes(n: Int): List[Int] =  
  primes(ints(2)) take n
```

## Note:

- `Stream.cons` lazily **prepends** the first argument to the second (a stream);
- `ints(n)` **lazily builds** an **infinite stream** of integers starting at `n`;
- `primes(ints)` lazily builds an infinite stream of prime numbers found in `ints`;
- `primes(n)` lazily builds a **finite stream** of the first `n` primes.

# Lazy Values

Lazy values **delay** computations, but do not run them concurrently.

```
val n = 1000
val ps = primes(n)
...
ps.foreach(/* process each prime */)
```

## Note:

- there is **no delay** when `ps` is assigned, and the thread proceeds;
- there is **delay** when `ps.foreach` is called;
- the whole computation of `primes(n)` is **enforced** before the first iteration of `foreach`;
- `val ps = primes(n)` is **functionally equivalent** to `lazy val ps = primes(n).toList`.

# Lazy Values

Lazy values **delay** computations, but do not run them concurrently.

```
val n = 1000
val ps = primes(n)
...
ps.foreach(/* process each prime */)
```

Question: How can we update the code so that:

- the computation of `primes(n)` actually **starts** when `ps` is assigned,
- the original thread **proceeds immediately** with further code,
- the thread is able to **check** whether computation of `ps` is complete, and then
- **if complete**, use `ps` without having to uselessly wait for it,
- otherwise, proceed with other tasks **instead of blocking**.

# Lazy Values

We can use shared state and wait for the child thread to complete:

```
val n = 1000
var ps: List[Int] = List.empty
val t = start { ps = primes(n) toList }
...
t.join()
ps.foreach(/* process each prime */)
```

**Note:** This approach is

- explicitly **waiting on thread completion**, not on the result;
- unable to **block conditionally**, only if the asynchronous computation is ready;
- error-prone: **shared state + concurrency = evil**



# Futures

```
val n = 1000
val future = Future { primes(n) }
...
while (!future.isCompleted) { /* proceed with other tasks while waiting */ }
future.value match {
  case Some(Success(ps)) => ps.foreach(/* process the prime */)
  case _ => }
```

## Note:

- `Future` is an **object**;
- `Future { primes(n) }` is short for `Future.apply(primes(n))`;
- `Future.apply` returns an instance of the **class** `Future`;
- `future` is a wrapper that will eventually be filled with a value;
- `future.isCompleted` is a **non-blocking** check of the future's state;
- `future.value` is the result of the computation.

# Futures

```
trait Future[+T] extends Awaitable[T] {  
  def value: Option[Try[T]]  
  ... }  
sealed abstract class Option[+T] ...  
case object None extends Option[Nothing] ...  
final case class Some[+T](x: T) extends Option[T] ...  
sealed abstract class Try[+T] ...  
final case class Success[+T](value: T) extends Try[T] ...  
final case class Failure[+T](exception: Throwable) extends Try[T] ...
```

## Note:

- `future.value` is an instance of `Option[Try[Stream[Int]]]`:
- until future is **completed**, `future.value` is `None`;
- if `primes(n)` **succeeded**, `future.value` is `Some(Success(ps))`;
- if `primes(n)` **failed** with exception `e`, `future.value` is `Some(Failure(e))`;
- `[+T]` denotes **type covariance**: if `U` is subclass of `T`, then `X[U]` is subclass of `X[T]` (think `<? extends T` in Java).

# Futures

Code depending on a future can be executed in a **callback**:

```
val future = Future { primes(n) }
future foreach {
  case ps => ps.foreach(p => log(s"$p")) }
future foreach {
  case ps => ps.zipWithIndex.foreach(ip => log(s"#{ip._2}: ${ip._1}") ) }
...
```

## Note:

- there are two separate **callbacks** installed on `future`;
- the two `future foreach` statements **return immediately**, irrespectively of `future.value`;
- the current thread proceeds without blocking;
- once `future` is complete, the callbacks can run **concurrently** to each other and the current thread.

# Futures

Callbacks can also handle failures—futures with `Failure[T]` as the value:

```
val future = Future { primes(n) }  
future foreach { case ps => /* process primes */ }  
(future failed) foreach { case e => /* process exception */ }  
...
```

## Note:

- `future foreach` **will only execute if** `future.value` is `Some(Success(ps))`;
- `(future failed) foreach` **will only execute if** `future.value` is `Some(Failure(e))`;
- `(future failed) foreach` is syntactic variant of `future.failed foreach` (and `future.failed.foreach`);
- `future onSuccess` and `future onFailure` are equivalent to the above;
- `future onComplete` installs a callback the will process `future` on both success and failure.

# Promises

Futures are

- **read-only** placeholders for values that **may not yet exist**.

Promises are

- **write-only, single-assignment** placeholders that
- **complete** futures using the `success` method, or
- **fail** futures using the `failure` method.

```
val promise: Promise[T] = Promise[T]
val future: Future[T] = promise.future
```

**Note:**

- a promise has an **associated future** object which
- can be used to **test** for and **read** the value that
- can be **set** through the promise object.

# Promises

```
val promise = Promise[T]
val future = promise future
future foreach { case result => /* handle result */ }
future.failed foreach { case error => /* handle failure */ }
start { promise success new T(/* provide value */) }
start { promise failure new Exception(/* provide failure message */) }
```

## Note:

- `Promise[T]` **constructs** a new promise object (for some type `T`);
- `promise future` is the **corresponding future** object;
- `promise success` **attempts** to assign the promise a **success value**;
- `promise failure` **attempts** to assign the promise a **failure exception**;
- **only one** of the above concurrent threads may succeed;
- when one did, the other **must fail**.

# Actors

The **actor model** of concurrent computing allows

- mutable state to be **entirely hidden** within individual actors;
- communication between actors to be **limited to passing messages**;
- **no memory** to be shared between actors;
- each actor to process its message queue **sequentially**; thus
- the order of processing of messages sent to different actors is **non-deterministic**,
- the order of processing of messages sent to the same actor is **deterministic**.

The Scala `Actors` library has been deprecated, the current default is the `akka` library—see [akka.io](http://akka.io).

# Actors

Consider a counter that can be used to keep track of passengers on a flight:

```
import akka.actor.Actor
class Counter(var count: Int = 0) extends Actor {
  override def receive: Receive = {
    case '+' => count += 1
    case '_' => count = 0
    case '.' => context.stop(self)
    case _ => } }
```

## Note:

- a `Counter` has an internal **mutable** state, initially set to 0;
- the method `receive` **matches a message** against a list of patterns;
- if it matches `'+'`, the counter increases its state by 1;
- if it matches `'_'`, the counter resets its state to 0;
- if it matches `'.'`, the counter **stops receiving** further messages;
- otherwise, the message is simply **ignored**.



# Actors

Actors are constructed using `Props`, `ActorRef` configuration objects:

```
val props: Props = Props(classOf[Counter], 0)
```

## Note:

- `classOf[T]` is the **runtime representation** of a class (think `T.class` in Java);
- `Props()` is short for `Props.apply()`;
- `Props(classOf[Counter], 0)` **configures** counters with initial state 0;
- providing 0 in `Props` is not necessary as `Counter` has a **default** for `count`.

# Actors

Actors are constructed using `Props`, `ActorRef` configuration objects:

```
val props: Props = Props(classOf[Counter], 0)
```

## Note:

- `classOf[T]` is the **runtime representation** of a class (think `T.class` in Java);
- `Props()` is short for `Props.apply()`;
- `Props(classOf[Counter], 0)` **configures** counters with initial state 0;
- providing 0 in `Props` is not necessary as `Counter` has a **default** for `count`.

It is customary to provide a **companion object** for the actor class:

```
object Counter {  
  def props = Props(new Counter)  
  def props(init: Int = 0): Props = Props(classOf[Counter], init) }  
}
```

- `props` is a `Props`-builder method.

# Actors

Actors are constructed and run within an `ActorSystem`:

```
val system: ActorSystem = ActorSystem("Counters")
val counter: ActorRef = system.actorOf(Counter.props, name="counter")
```

## Note:

- `ActorSystem()` is short for `ActorSystem.apply()`;
- `system` is an `ActorSystem` with the name `"Counters"`;
- `Counter.props` creates a `Props` object;
- `counter` is an `ActorRef`, a **reference to an actor** within `system`;
- `counter` has the name `"counter"` within `system`.

# Actors

Actors receive and process messages:

```
counter ! '+' // 1  
counter ! '+' // 2  
counter ! '_' // 3  
counter ! '?' // 4  
counter ! '.' // 5  
counter ! '+' // 6
```

# Actors

Actors receive and process messages:

```
counter ! '+' // 1
counter ! '+' // 2
counter ! '_' // 3
counter ! '?' // 4
counter ! '.' // 5
counter ! '+' // 6
```

## Note:

- 1 and 2 bump counter to 1, then 2;
- 3 resets counter to 0;
- 4 sends an unrecognized message;
- 5 stops counter from receiving further messages;
- 6 is never delivered.

# Actors

Actors can handle messages differently depending on their state:

```
class Counter(var count: Int = 0) extends Actor {  
  def receive = unlocked  
  def unlocked: Receive = {  
    case '+' => count += 1  
    ...  
    case '*' => context.become(locked) }  
  def locked: Receive = {  
    case '*' => context.become(unlocked)  
    case _ => } }  
}
```

## Note:

- initially, `receive` is `unlocked` and the counter acts as previously;
- if the message is `'*'`, `receive` is changed to `locked`;
- `locked` ignores all messages except for `'*'`, on which it restores `receive` to `unlocked`.

# Actors

Actors can **send responses** to their senders:

```
class Counter(var count: Int = 0) extends Actor {  
  def receive = {  
    case '=' => sender ! count  
    ... } }  
...  
val response = counter ? '='  
response match { case Some(Success(count)) => /* handle count */ }
```

## Note:

- `counter ? '='` sends an **ask** request to `counter`;
- in `receive`, if the message is `=`, the actor sends `count` to the **sender** of the message;
- `sender` is a **method inherited** from `Actor`;
- `response` is a **future**;
- if the operation succeeded, `response` contains `Some(Success(count))`.